

Team 2 - Homework Two

Assignment 2: KJ 7.2; KJ 7.5

Sang Yoon (Andy) Hwang

DATE:2019-11-01

Dependencies

```
# predictive modeling
libraries('mlbench', 'caret', 'AppliedPredictiveModeling')

# Formatting Libraries
libraries('default', 'knitr', 'kableExtra')

# Plotting Libraries
libraries('ggplot2', 'grid', 'ggfortify')
```

(1) Kuhn & Johnson 7.2

Friedman (1991) introduced several benchmark data sets create by simulation. One of these simulations used the following nonlinear equation to create data: $y = 10\sin(\pi x_1 x_2)20(x_3 - 0.5)^2 10x_4 5x_5 N(0, \sigma^2)$; where the x values are random variables uniformly distributed between $[0, 1]$ (there are also 5 other non-informative variables also created in the simulation).


****The package `mlbench` contains a function called `mlbench.friedman1` that simulates these data:****

```
set.seed(200)
trainingData <- mlbench.friedman1(200, sd = 1)

## We convert the 'x' data from a matrix to a data frame
## One reason is that this will give the columns names.

trainingData$x <- data.frame(trainingData$x)

## Look at the data using
featurePlot(trainingData$x, trainingData$y)
```



```
## or other methods.
```

```
## This creates a list with a vector 'y' and a matrix
```

```
## of predictors 'x'. Also simulate a large test set to
## estimate the true error rate with good precision:
```

```
testData <- mlbench.friedman1(5000, sd = 1)
testData$x <- data.frame(testData$x)
```

(a) Tune several models on these data. For example:

```
set.seed(200)
knnModel <- train(x = trainingData$x,
                  y = trainingData$y,
                  method = "knn",
                  preProc = c("center", "scale"),
                  tuneLength = 10)
knnModel
```

k-Nearest Neighbors

200 samples
10 predictor

Pre-processing: centered (10), scaled (10)

Resampling: Bootstrapped (25 reps)

Summary of sample sizes: 200, 200, 200, 200, 200, 200, ...

Resampling results across tuning parameters:

k	RMSE	Rsquared	MAE
5	3.554554	0.4895311	2.883670
7	3.423303	0.5264402	2.767070
9	3.361439	0.5525056	2.702852
11	3.275234	0.5885952	2.632363
13	3.245376	0.6099949	2.607849
15	3.218637	0.6308597	2.576730
17	3.229692	0.6380326	2.589076
19	3.231915	0.6463092	2.595749
21	3.228217	0.6591640	2.599618
23	3.254794	0.6610119	2.628392

RMSE was used to select the optimal model using the smallest value.

The final value used for the model was k = 15.

```
knnPred <- predict(knnModel, newdata = testData$x)
```

```
## The function 'postResample' can be used to get the test set performance values
postResample(pred = knnPred, obs = testData$y)
```

```
      RMSE  Rsquared      MAE
3.1750657 0.6785946 2.5443169
```

Model 1: KNN model with hyperparameter tuning

```
set.seed(100)
knn_model <- train(trainingData$x,
                  trainingData$y,
                  method = "knn",
                  # Center and scaling will occur for new predictions too
```

```
preProc = c("center", "scale"),
tuneGrid = data.frame(.k = 1:50),
trControl = trainControl(method = "cv"))
```

knn_model

k-Nearest Neighbors

200 samples

10 predictor

Pre-processing: centered (10), scaled (10)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...

Resampling results across tuning parameters:

k	RMSE	Rsquared	MAE
1	4.218081	0.4124441	3.500924
2	3.546791	0.5197988	2.941632
3	3.460792	0.5170308	2.841547
4	3.304737	0.5584950	2.697459
5	3.281638	0.5736341	2.695889
6	3.201700	0.6056758	2.651120
7	3.212544	0.6083203	2.677007
8	3.175321	0.6293337	2.589611
9	3.117499	0.6632514	2.555966
10	3.118730	0.6619864	2.542520
11	3.059105	0.6860779	2.486276
12	3.110398	0.6794240	2.539603
13	3.107626	0.6909955	2.504201
14	3.098502	0.6981615	2.503473
15	3.107763	0.7020139	2.517724
16	3.092651	0.7141572	2.499280
17	3.094870	0.7184612	2.502765
18	3.106069	0.7205098	2.512843
19	3.110156	0.7223952	2.522663
20	3.129738	0.7248629	2.544792
21	3.125325	0.7282228	2.555344
22	3.152564	0.7265154	2.583124
23	3.156426	0.7299184	2.590130
24	3.156312	0.7315798	2.583509
25	3.168639	0.7320840	2.600023
26	3.192311	0.7301789	2.615933
27	3.191953	0.7359772	2.626247
28	3.236432	0.7229397	2.667947
29	3.266908	0.7199837	2.696586
30	3.264200	0.7279598	2.695916
31	3.273003	0.7273885	2.702978
32	3.282489	0.7276198	2.702930
33	3.289932	0.7284479	2.707697
34	3.303444	0.7281557	2.717839
35	3.320845	0.7249744	2.727184
36	3.331038	0.7257760	2.729478
37	3.342847	0.7296750	2.740671

38	3.353494	0.7305494	2.745450
39	3.363513	0.7335993	2.748193
40	3.369217	0.7349916	2.752439
41	3.394484	0.7291182	2.779272
42	3.404124	0.7292558	2.783362
43	3.403433	0.7342789	2.780712
44	3.418392	0.7341581	2.794006
45	3.420892	0.7410895	2.790850
46	3.421315	0.7485185	2.801475
47	3.434838	0.7498270	2.816174
48	3.449996	0.7487060	2.831474
49	3.451778	0.7525105	2.833605
50	3.467105	0.7471964	2.844317

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was k = 11.

```
knn_Pred <- predict(knn_model, newdata = testData$x)

## The function 'postResample' can be used to get the test set performance values
knn_pv <- postResample(pred = knn_Pred, obs = testData$y)
knn_pv
```

	RMSE	Rsquared	MAE
	3.1222641	0.6690472	2.4963650

Unlike above approach where `tuneLength = 10` to find 10 odd numbered Ks starting from 5, we will set `tuneGrid` running from k = 1 to 50 after CV process. RMSE on validation set was used to select the optimal model using the smallest value. The final value used for the model was k = 11 with RMSE on test set of 3.1222641.

Model 2: Neural Networks

```
# remove highly correlated predictors to ensure that the maximum absolute pairwise correlation between
# we did not have any highly correlated predictors so let's keep the features as they are.
findCorrelation(cor(trainingData$x), cutoff = .75)
```

```
integer(0)

# hyperparameter tuning for nnet
nnetGrid <- expand.grid(.size = c(1:10), .decay = c(0, 0.01, .1))

set.seed(100)
nnet_model <- train(trainingData$x, trainingData$y,
  method = "nnet",
  tuneGrid = nnetGrid,
  trControl = trainControl(method="cv"),
  ## Automatically standardize data prior to modeling and prediction
  preProc = c("center", "scale"),
  linout = TRUE,
  trace = FALSE,
  MaxNWts = 10 * (ncol(trainingData$x) + 1) + 10 + 1,
  maxit = 500)

nnet_model
```

Neural Network

200 samples
10 predictor

Pre-processing: centered (10), scaled (10)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...

Resampling results across tuning parameters:

size	decay	RMSE	Rsquared	MAE
1	0.00	2.803056	0.6666327	2.258691
1	0.01	2.427596	0.7621743	1.887808
1	0.10	2.435471	0.7608686	1.890066
2	0.00	2.489139	0.7430098	1.959395
2	0.01	2.580965	0.7303017	2.026386
2	0.10	2.697160	0.7081857	2.145994
3	0.00	2.175481	0.8114073	1.736758
3	0.01	2.247174	0.7923587	1.847530
3	0.10	2.574029	0.7395256	2.056691
4	0.00	2.338196	0.7863493	1.857565
4	0.01	2.382258	0.7789730	1.875869
4	0.10	2.441937	0.7640421	1.906016
5	0.00	4.095639	0.6461716	2.815675
5	0.01	2.611351	0.7275315	2.058338
5	0.10	2.527299	0.7479751	2.043740
6	0.00	4.360771	0.5655777	2.776134
6	0.01	2.740151	0.7334422	2.120769
6	0.10	2.723467	0.7037555	2.191856
7	0.00	8.107456	0.6013541	3.614263
7	0.01	2.634878	0.7306846	2.082135
7	0.10	2.634554	0.7272090	2.034430
8	0.00	9.294033	0.4764955	3.995120
8	0.01	3.239373	0.6734412	2.543492
8	0.10	2.857949	0.7043441	2.285976
9	0.00	7.318398	0.5467571	3.670600
9	0.01	3.541729	0.5379448	2.671080
9	0.10	3.187206	0.6050659	2.544700
10	0.00	4.083062	0.5181907	2.838082
10	0.01	3.525950	0.5879990	2.836265
10	0.10	3.010387	0.6719106	2.387353

RMSE was used to select the optimal model using the smallest value.

The final values used for the model were size = 3 and decay = 0.

```
nnet_Pred <- predict(nnet_model, newdata = testData$x)
```

```
## The function 'postResample' can be used to get the test set performance values
```

```
nnet_pv <- postResample(pred = nnet_Pred, obs = testData$y)
```

```
nnet_pv
```

RMSE	Rsquared	MAE
2.3950120	0.7740742	1.7970761

We found nnet with size = 3 (number of units in the hidden layer) and decay = 0 (parameter for weight decay) is the optimal model based on RMSE on validating set. RMSE on test set was 2.395012.

Model 3: Neural Networks Using Model Averaging

```
# hyperparameter tuning for avnnet
nnetGrid2 <- expand.grid(.size = c(1:10), .decay = c(0, 0.01, .1), .bag = FALSE)

set.seed(100)
avnnet_model <- train(trainingData$x, trainingData$y,
  method = "avNNet",
  tuneGrid = nnetGrid2,
  trControl = trainControl(method="cv"),
  ## Automatically standardize data prior to modeling and prediction
  preProc = c("center", "scale"),
  linout = TRUE,
  trace = FALSE,
  MaxNWts = 10 * (ncol(trainingData$x) + 1) + 10 + 1,
  maxit = 500)

avnnet_model
```

Model Averaged Neural Network

200 samples
10 predictor

Pre-processing: centered (10), scaled (10)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...

Resampling results across tuning parameters:

size	decay	RMSE	Rsquared	MAE
1	0.00	2.451722	0.7625351	1.896021
1	0.01	2.427658	0.7621763	1.887707
1	0.10	2.435387	0.7608838	1.890037
2	0.00	2.458303	0.7556215	1.947825
2	0.01	2.447743	0.7533484	1.884509
2	0.10	2.483261	0.7506467	1.924237
3	0.00	2.086103	0.8297667	1.625836
3	0.01	2.155718	0.8096710	1.676070
3	0.10	2.227768	0.8003473	1.734663
4	0.00	2.036944	0.8294847	1.636994
4	0.01	2.102926	0.8232439	1.623997
4	0.10	2.032960	0.8301759	1.604834
5	0.00	2.314361	0.7871277	1.717315
5	0.01	2.131569	0.8131597	1.670786
5	0.10	2.192707	0.8054839	1.778156
6	0.00	2.593804	0.7448382	1.966474
6	0.01	2.167612	0.8108275	1.674682
6	0.10	2.042238	0.8227937	1.645410
7	0.00	4.989121	0.5232150	3.110796
7	0.01	2.238776	0.7948069	1.754542
7	0.10	2.202599	0.7926956	1.751645
8	0.00	5.839958	0.5788055	3.266803
8	0.01	2.392621	0.7727856	1.893142
8	0.10	2.271702	0.7860966	1.790783
9	0.00	5.138515	0.4191673	3.248740

9	0.01	2.463813	0.7534985	1.927194
9	0.10	2.250755	0.7930757	1.820346
10	0.00	3.400511	0.6426513	2.428742
10	0.01	2.456849	0.7404324	1.988814
10	0.10	2.437919	0.7608963	1.971687

Tuning parameter 'bag' was held constant at a value of FALSE
 RMSE was used to select the optimal model using the smallest value.
 The final values used for the model were size = 4, decay = 0.1 and bag
 = FALSE.

```
avnnet_Pred <- predict(avnnet_model, newdata = testData$x)

## The function 'postResample' can be used to get the test set performance values
avnnet_pv <- postResample(pred = avnnet_Pred, obs = testData$y)
avnnet_pv
```

RMSE	Rsquared	MAE
2.1306481	0.8202697	1.5982639

We found nnet with size = 4 (number of units in the hidden layer) and decay = 0.1 (parameter for weight decay) is the optimal model based on RMSE on validating set. RMSE on test set was 2.1306481.

Model 4: Multivariate Adaptive Regression Splines

```
# hyperparameter tuning for MARS
marsGrid <- expand.grid(.degree = 1:3, .nprune = 2:38)

set.seed(100)
mars_model <- train(trainingData$x, trainingData$y,
  method = "earth",
  tuneGrid = marsGrid,
  trControl = trainControl(method="cv"))

#mars_model
summary(mars_model)
```

Call: earth(x=data.frame[200,10], y=c(18.46,16.1,17...), keepxy=TRUE,
 degree=2, nprune=17)

	coefficients
(Intercept)	20.378441
h(0.621722-X1)	-15.512132
h(X1-0.621722)	9.177132
h(0.601063-X2)	-17.940676
h(X2-0.601063)	10.064356
h(X3-0.281766)	11.590022
h(0.447442-X3)	14.641640
h(X3-0.447442)	-12.924806
h(X3-0.606015)	13.416764
h(0.734892-X4)	-10.074386
h(X4-0.734892)	9.687149
h(0.850094-X5)	-5.385762
h(0.218266-X1) * h(X2-0.601063)	-55.372637
h(X1-0.218266) * h(X2-0.601063)	-27.542831
h(X1-0.621722) * h(X2-0.295997)	-26.527403

```
h(0.649253-X1) * h(0.601063-X2)    26.129827
```

Selected 16 of 18 terms, and 5 of 10 predictors

Termination condition: Reached nk 21

Importance: X1, X4, X2, X5, X3, X6-unused, X7-unused, X8-unused, ...

Number of terms at each degree of interaction: 1 11 4

GCV 1.61518 RSS 210.6377 GRSq 0.934423 RSq 0.9568093

```
mars_Pred <- predict(mars_model, newdata = testData$x)
```

```
## The function 'postResample' can be used to get the test set performance values
```

```
mars_pv <- postResample(pred = mars_Pred, obs = testData$y)
```

```
mars_pv
```

```
      RMSE Rsquared      MAE
1.1492504 0.9471145 0.9158382
```

We found MARS with `degree = 2` (Maximum degree of interaction (Friedman's *mi*)) and `nprune = 17` (maximum number of terms (including intercept) in the pruned model) is the optimal model based on RMSE on validating set. RMSE on test set was 1.1492504.

Model 5: Support Vector regression

```
set.seed(100)
svm_model <- train(trainingData$x, trainingData$y,
  method = "svmRadial",
  preProc = c("center", "scale"),
  tuneLength = 14,
  trControl = trainControl(method="cv"))
```

```
#svm_model
```

```
svm_model
```

Support Vector Machines with Radial Basis Function Kernel

200 samples

10 predictor

Pre-processing: centered (10), scaled (10)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...

Resampling results across tuning parameters:

C	RMSE	Rsquared	MAE
0.25	2.534788	0.7882081	2.034824
0.50	2.292127	0.8029516	1.819981
1.00	2.091598	0.8284381	1.657402
2.00	1.967193	0.8457471	1.546737
4.00	1.883133	0.8561761	1.482054
8.00	1.863807	0.8588797	1.468328
16.00	1.834215	0.8633819	1.456738
32.00	1.836471	0.8632508	1.459909
64.00	1.836471	0.8632508	1.459909
128.00	1.836471	0.8632508	1.459909
256.00	1.836471	0.8632508	1.459909
512.00	1.836471	0.8632508	1.459909


```
1024.00 1.836471 0.8632508 1.459909
2048.00 1.836471 0.8632508 1.459909
```

Tuning parameter 'sigma' was held constant at a value of 0.0552698
 RMSE was used to select the optimal model using the smallest value.
 The final values used for the model were sigma = 0.0552698 and C = 16.

```
svm_Pred <- predict(svm_model, newdata = testData$x)

## The function 'postResample' can be used to get the test set performance values
svm_pv <- postResample(pred = svm_Pred, obs = testData$y)
svm_pv
```

```
      RMSE Rsquared      MAE
2.0490047 0.8297577 1.5586106
```

Since the nature of the equation of the data is non-linear, we will use `svmRadial` as kernel function for regression. The final values used for the model were sigma = 0.0552698 and C = 16 with RMSE on test set of 2.0490047.

- (b) Which models appear to give the best performance? Does MARS select the informative predictors (those named X1-X5)?

MARS appears to give the best performance based on RMSE, R squared and MAE on test set. The summary output of `mars_model` gives us that Importance: X1, X4, X2, X5, X3, X6-unused, X7-unused, X8-unused, X9-unused, MARS does select the informative predictors X1-X5 only.

```
#code
# Model performance metrics
sum_t <- data.frame(
  knn_pv,
  nnet_pv,
  avnnet_pv,
  mars_pv,
  svm_pv
)
print(sum_t)

      knn_pv  nnet_pv avnnet_pv  mars_pv  svm_pv
RMSE      3.1222641 2.3950120 2.1306481 1.1492504 2.0490047
Rsquared  0.6690472 0.7740742 0.8202697 0.9471145 0.8297577
MAE       2.4963650 1.7970761 1.5982639 0.9158382 1.5586106

# summary mars
summary(mars_model)
```

```
Call: earth(x=data.frame[200,10], y=c(18.46,16.1,17...), keepxy=TRUE,
  degree=2, nprune=17)
```

```
coefficients
(Intercept)      20.378441
h(0.621722-X1)   -15.512132
h(X1-0.621722)    9.177132
h(0.601063-X2)   -17.940676
h(X2-0.601063)   10.064356
h(X3-0.281766)   11.590022
h(0.447442-X3)   14.641640
h(X3-0.447442)  -12.924806
```

```

h(X3-0.606015)          13.416764
h(0.734892-X4)         -10.074386
h(X4-0.734892)          9.687149
h(0.850094-X5)         -5.385762
h(0.218266-X1) * h(X2-0.601063) -55.372637
h(X1-0.218266) * h(X2-0.601063) -27.542831
h(X1-0.621722) * h(X2-0.295997) -26.527403
h(0.649253-X1) * h(0.601063-X2)  26.129827

```

Selected 16 of 18 terms, and 5 of 10 predictors

Termination condition: Reached nk 21

Importance: X1, X4, X2, X5, X3, X6-unused, X7-unused, X8-unused, ...

Number of terms at each degree of interaction: 1 11 4

GCV 1.61518 RSS 210.6377 GRSq 0.934423 RSq 0.9568093

(2) Kuhn & Johnson 7.5

Exercise 6.3 describes data for a chemical manufacturing process. Use the same data imputation, data splitting, and pre-processing steps as before and train several nonlinear regression models.

- (a) Which nonlinear regression model gives the optimal resampling and test set performance?

```

# code
# split data train/test
training <- df_final$Yield %>%
  createDataPartition(p = 0.8, list = FALSE)

df_train <- df_final[training, ]
df_test <- df_final[-training, ]

# model1 - KNN
set.seed(100)
knn_model2 <- train(Yield~., data = df_train,
  method = "knn",
  # Center and scaling will occur for new predictions too
  preProc = c("center", "scale"),
  tuneGrid = data.frame(.k = 1:50),
  trControl = trainControl(method = "cv"))

knn_model2

```

k-Nearest Neighbors

144 samples

56 predictor

Pre-processing: centered (56), scaled (56)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 129, 131, 128, 129, 129, 130, ...

Resampling results across tuning parameters:

k	RMSE	Rsquared	MAE
1	1.562947	0.4104086	1.211757
2	1.394114	0.4526620	1.103294

3	1.285798	0.5202328	1.022250
4	1.261951	0.5511694	1.005031
5	1.290154	0.5346045	1.016046
6	1.285469	0.5467205	1.015962
7	1.277158	0.5691222	1.014170
8	1.280950	0.5672804	1.022692
9	1.260191	0.5951063	1.008777
10	1.268148	0.5927028	1.022267
11	1.274834	0.5762849	1.040308
12	1.283879	0.5691015	1.046630
13	1.300146	0.5565749	1.063440
14	1.294960	0.5643696	1.054302
15	1.315419	0.5533099	1.071270
16	1.322033	0.5432361	1.078412
17	1.333203	0.5374531	1.078099
18	1.345231	0.5267062	1.089399
19	1.354941	0.5240723	1.095199
20	1.361053	0.5260280	1.100479
21	1.357930	0.5291765	1.095387
22	1.360462	0.5363972	1.100225
23	1.369994	0.5307997	1.105951
24	1.381541	0.5262263	1.116182
25	1.390061	0.5225105	1.124752
26	1.390820	0.5240722	1.125271
27	1.401032	0.5200972	1.133713
28	1.406645	0.5247034	1.137767
29	1.417484	0.5124727	1.145324
30	1.422851	0.5160566	1.149150
31	1.424039	0.5158770	1.155165
32	1.431312	0.5132871	1.164555
33	1.441439	0.5029494	1.170119
34	1.445577	0.5033464	1.173391
35	1.446068	0.4997379	1.173717
36	1.451050	0.4987082	1.175739
37	1.450722	0.5073898	1.176310
38	1.452006	0.5093112	1.180577
39	1.458957	0.5080230	1.185948
40	1.462371	0.5046533	1.190404
41	1.470614	0.5009788	1.197493
42	1.476481	0.4945194	1.201055
43	1.478755	0.4983923	1.199669
44	1.484498	0.4951622	1.203632
45	1.490334	0.4982518	1.208023
46	1.491714	0.5057906	1.210398
47	1.496813	0.5022176	1.215699
48	1.500448	0.5047809	1.219058
49	1.506809	0.5048347	1.225709
50	1.509769	0.5034198	1.228911

RMSE was used to select the optimal model using the smallest value.
The final value used for the model was k = 9.

```
knn_Pred2 <- predict(knn_model2, newdata = df_test)
```

```
## The function 'postResample' can be used to get the test set performance values
knn_pv2 <- postResample(pred = knn_Pred2, obs = df_test$Yield)

# model2 - nnet
# remove highly correlated predictors to ensure that the maximum absolute pairwise correlation between
df_train_x <- df_train[-1]
df_train_y <- df_train[,1]
df_test_x <- df_test[-1]
df_test_y <- df_test[,1]

tooHigh <- findCorrelation(cor(df_train_x), cutoff = .75)

trainx_nn <- df_train_x[, -tooHigh]
testx_nn <- df_test_x[, -tooHigh]

# hyperparameter tuning for nnet
nnetGrid12 <- expand.grid(.size = c(1:10), .decay = c(0, 0.01, .1))

set.seed(100)
nnet_model2 <- train(trainx_nn, df_train_y,
  method = "nnet",
  tuneGrid = nnetGrid12,
  trControl = trainControl(method="cv"),
  ## Automatically standardize data prior to modeling and prediction
  preProc = c("center", "scale"),
  linout = TRUE,
  trace = FALSE,
  MaxNWts = 10 * (ncol(trainx_nn) + 1) + 10 + 1,
  maxit = 500)

nnet_model2
```

Neural Network

144 samples

35 predictor

Pre-processing: centered (35), scaled (35)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 129, 131, 128, 129, 129, 130, ...

Resampling results across tuning parameters:

size	decay	RMSE	Rsquared	MAE
1	0.00	1.755205	0.2066616	1.455577
1	0.01	1.531002	0.4448691	1.192716
1	0.10	1.373454	0.5175944	1.109371
2	0.00	2.212948	0.1773784	1.727733
2	0.01	1.864348	0.3910941	1.542705
2	0.10	1.715287	0.4696443	1.333888
3	0.00	4.213947	0.2706156	3.025339
3	0.01	3.402447	0.1312987	2.650608
3	0.10	2.274638	0.3535430	1.809244
4	0.00	3.753927	0.1694120	2.927575

4	0.01	3.032265	0.2730534	2.293653
4	0.10	2.662301	0.2596654	2.073559
5	0.00	3.193235	0.2142524	2.523091
5	0.01	2.847820	0.2693243	2.167825
5	0.10	2.843172	0.2010863	2.136950
6	0.00	3.141860	0.3508280	2.459695
6	0.01	2.535074	0.1739787	2.089114
6	0.10	2.009758	0.3223722	1.549843
7	0.00	5.113950	0.1935983	3.377113
7	0.01	2.849091	0.2545805	2.129850
7	0.10	2.114057	0.2762636	1.594852
8	0.00	4.294256	0.1767868	3.299536
8	0.01	2.795623	0.1388297	2.169489
8	0.10	2.275122	0.2983466	1.713263
9	0.00	9.006573	0.2327137	5.799823
9	0.01	2.730389	0.2517318	2.214361
9	0.10	2.237043	0.3159019	1.714727
10	0.00	8.689726	0.3538469	5.208147
10	0.01	4.329962	0.1650993	3.119468
10	0.10	2.410800	0.2083370	1.808763

RMSE was used to select the optimal model using the smallest value.
The final values used for the model were size = 1 and decay = 0.1.

```
nnet_Pred2 <- predict(nnet_model2, newdata = testx_nn)

## The function 'postResample' can be used to get the test set performance values
nnet_pv2 <- postResample(pred = nnet_Pred2, obs = df_test_y)

# model 3 - avNNet
# hyperparameter tuning for avnnet
nnetGrid22 <- expand.grid(.size = c(1:10), .decay = c(0, 0.01, .1), .bag = FALSE)

set.seed(100)
avnnet_model2 <- train(trainx_nn, df_train_y,
  method = "avNNet",
  tuneGrid = nnetGrid22,
  trControl = trainControl(method="cv"),
  ## Automatically standardize data prior to modeling and prediction
  preProc = c("center", "scale"),
  linout = TRUE,
  trace = FALSE,
  MaxNWts = 10 * (ncol(trainx_nn) + 1) + 10 + 1,
  maxit = 500)

avnnet_model2
```

Model Averaged Neural Network

144 samples
35 predictor

Pre-processing: centered (35), scaled (35)

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 129, 131, 128, 129, 129, 130, ...

Resampling results across tuning parameters:

size	decay	RMSE	Rsquared	MAE
1	0.00	1.589451	0.3851410	1.271954
1	0.01	1.376260	0.5016400	1.103760
1	0.10	1.298216	0.5536547	1.053173
2	0.00	1.614710	0.4230801	1.274696
2	0.01	1.435679	0.5191818	1.116245
2	0.10	1.486560	0.4888622	1.185787
3	0.00	1.893573	0.3390867	1.482814
3	0.01	1.947427	0.2967394	1.543828
3	0.10	1.702341	0.4794429	1.298251
4	0.00	2.183278	0.2198895	1.802826
4	0.01	1.837276	0.3676411	1.438189
4	0.10	1.841018	0.3934255	1.451878
5	0.00	1.788868	0.3797231	1.487272
5	0.01	1.677605	0.4073755	1.350633
5	0.10	1.790401	0.4160895	1.340008
6	0.00	2.183750	0.3741136	1.680205
6	0.01	1.474744	0.5088465	1.228068
6	0.10	1.842838	0.4122807	1.340819
7	0.00	2.535701	0.3775290	2.006750
7	0.01	1.678518	0.4521095	1.327766
7	0.10	1.808841	0.3698767	1.370640
8	0.00	3.327318	0.2290038	2.309492
8	0.01	1.619979	0.4401995	1.268287
8	0.10	1.789302	0.4025716	1.399899
9	0.00	4.247876	0.2434059	2.861125
9	0.01	1.699004	0.4937072	1.303050
9	0.10	1.967303	0.2727353	1.492640
10	0.00	7.011010	0.2117344	4.342906
10	0.01	2.038958	0.3644072	1.468109
10	0.10	1.748484	0.3668079	1.353886

Tuning parameter 'bag' was held constant at a value of FALSE

RMSE was used to select the optimal model using the smallest value.

The final values used for the model were size = 1, decay = 0.1 and bag = FALSE.

```
avnnet_Pred2 <- predict(avnnet_model2, newdata = testx_nn)

## The function 'postResample' can be used to get the test set performance values
avnnet_pv2 <- postResample(pred = avnnet_Pred2, obs = df_test_y)

# model 4 - MARS
# hyperparameter tuning for MARS
marsGrid2 <- expand.grid(.degree = 1:3, .nprune = 2:38)

set.seed(100)
mars_model2 <- train(Yield~., data = df_train,
                     method = "earth",
                     tuneGrid = marsGrid2,
                     trControl = trainControl(method="cv"))
```

```
#mars_model
mars_model2
```

Multivariate Adaptive Regression Spline

144 samples

56 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 129, 131, 128, 129, 129, 130, ...

Resampling results across tuning parameters:

degree	nprune	RMSE	Rsquared	MAE
1	2	1.407400	0.4656139	1.1174779
1	3	1.223427	0.6194501	0.9868635
1	4	1.199954	0.6230092	0.9636997
1	5	1.164765	0.6453235	0.9463908
1	6	1.150927	0.6495515	0.9307132
1	7	1.098522	0.6828422	0.8871423
1	8	1.150254	0.6602017	0.9190847
1	9	1.163031	0.6570934	0.9352294
1	10	1.202534	0.6312059	0.9743665
1	11	1.185338	0.6391984	0.9467194
1	12	1.176559	0.6346864	0.9469068
1	13	1.211428	0.6172954	0.9842769
1	14	1.217020	0.6188990	0.9886593
1	15	1.223154	0.6160753	0.9926387
1	16	1.223154	0.6160753	0.9926387
1	17	1.223154	0.6160753	0.9926387
1	18	1.223154	0.6160753	0.9926387
1	19	1.223154	0.6160753	0.9926387
1	20	1.223154	0.6160753	0.9926387
1	21	1.223154	0.6160753	0.9926387
1	22	1.223154	0.6160753	0.9926387
1	23	1.223154	0.6160753	0.9926387
1	24	1.223154	0.6160753	0.9926387
1	25	1.223154	0.6160753	0.9926387
1	26	1.223154	0.6160753	0.9926387
1	27	1.223154	0.6160753	0.9926387
1	28	1.223154	0.6160753	0.9926387
1	29	1.223154	0.6160753	0.9926387
1	30	1.223154	0.6160753	0.9926387
1	31	1.223154	0.6160753	0.9926387
1	32	1.223154	0.6160753	0.9926387
1	33	1.223154	0.6160753	0.9926387
1	34	1.223154	0.6160753	0.9926387
1	35	1.223154	0.6160753	0.9926387
1	36	1.223154	0.6160753	0.9926387
1	37	1.223154	0.6160753	0.9926387
1	38	1.223154	0.6160753	0.9926387
2	2	1.407400	0.4656139	1.1174779
2	3	1.280199	0.5709792	1.0181681
2	4	1.251083	0.5829880	1.0166939

2	5	1.194031	0.6200359	0.9658520
2	6	1.179059	0.6219074	0.9538453
2	7	1.091866	0.6786885	0.8650013
2	8	1.144683	0.6540160	0.9054151
2	9	1.199698	0.6234861	0.9292222
2	10	1.207126	0.6211522	0.9260997
2	11	1.278838	0.6183028	0.9523748
2	12	1.304635	0.6120710	0.9668832
2	13	1.339764	0.6048324	0.9990698
2	14	1.410092	0.5885196	1.0481210
2	15	1.411828	0.5860411	1.0611942
2	16	1.750442	0.5340995	1.1622914
2	17	1.786237	0.5342503	1.1606896
2	18	1.781165	0.5339347	1.1559614
2	19	1.783391	0.5384735	1.1460005
2	20	1.796314	0.5288854	1.1652398
2	21	2.062368	0.5207735	1.2547404
2	22	2.088379	0.5172560	1.2679488
2	23	2.066402	0.5309589	1.2665940
2	24	2.065437	0.5344523	1.2633582
2	25	2.079735	0.5316228	1.2667648
2	26	2.081797	0.5324254	1.2677712
2	27	2.081797	0.5324254	1.2677712
2	28	2.081797	0.5324254	1.2677712
2	29	2.081797	0.5324254	1.2677712
2	30	2.081797	0.5324254	1.2677712
2	31	2.081797	0.5324254	1.2677712
2	32	2.081797	0.5324254	1.2677712
2	33	2.081797	0.5324254	1.2677712
2	34	2.081797	0.5324254	1.2677712
2	35	2.081797	0.5324254	1.2677712
2	36	2.081797	0.5324254	1.2677712
2	37	2.081797	0.5324254	1.2677712
2	38	2.081797	0.5324254	1.2677712
3	2	1.407400	0.4656139	1.1174779
3	3	1.288493	0.5682931	1.0263117
3	4	1.379392	0.5892361	1.0481819
3	5	1.366985	0.5842811	1.0532546
3	6	1.363445	0.5710827	1.0519441
3	7	1.317884	0.6110113	0.9876128
3	8	2.155901	0.5337454	1.2717478
3	9	2.365837	0.5628087	1.3261618
3	10	1.345061	0.6164859	1.0209735
3	11	1.450749	0.5683842	1.0699845
3	12	1.436353	0.5685230	1.0624539
3	13	1.410479	0.5816416	1.0382527
3	14	1.715407	0.5519921	1.1549908
3	15	2.828992	0.4688554	1.4964577
3	16	2.831170	0.4265037	1.5243968
3	17	2.780085	0.4308703	1.5035084
3	18	2.807375	0.4281497	1.4931957
3	19	3.065332	0.4132803	1.6030211
3	20	3.042777	0.4181071	1.5828211
3	21	3.069110	0.4190439	1.5796859

3	22	2.958982	0.4175394	1.5496401
3	23	2.955995	0.4169054	1.5485249
3	24	2.973481	0.4108560	1.5689330
3	25	3.019521	0.3880575	1.5958254
3	26	3.024177	0.3822188	1.5964017
3	27	2.878881	0.3835998	1.5462256
3	28	3.013091	0.3826830	1.5781993
3	29	3.013091	0.3826830	1.5781993
3	30	3.013091	0.3826830	1.5781993
3	31	3.013091	0.3826830	1.5781993
3	32	3.013091	0.3826830	1.5781993
3	33	3.013091	0.3826830	1.5781993
3	34	3.013091	0.3826830	1.5781993
3	35	3.013091	0.3826830	1.5781993
3	36	3.013091	0.3826830	1.5781993
3	37	3.013091	0.3826830	1.5781993
3	38	3.013091	0.3826830	1.5781993

RMSE was used to select the optimal model using the smallest value.
The final values used for the model were `nprune = 7` and `degree = 2`.

```

mars_Pred2 <- predict(mars_model2, newdata = df_test)

## The function 'postResample' can be used to get the test set performance values
mars_pv2 <- postResample(pred = mars_Pred2, obs = df_test$Yield)

# model 5 - SVM - regression
set.seed(100)
svm_model2 <- train(Yield~., data = df_train,
  method = "svmRadial",
  preProc = c("center", "scale"),
  tuneLength = 14,
  trControl = trainControl(method="cv"))

#svm_model
svm_model2

```

Support Vector Machines with Radial Basis Function Kernel

144 samples
56 predictor

Pre-processing: centered (56), scaled (56)
Resampling: Cross-Validated (10 fold)
Summary of sample sizes: 129, 131, 128, 129, 129, 130, ...
Resampling results across tuning parameters:

C	RMSE	Rsquared	MAE
0.25	1.372028	0.5607359	1.1092751
0.50	1.237001	0.6179908	1.0047103
1.00	1.126298	0.6748623	0.9158092
2.00	1.081110	0.6935348	0.8707261
4.00	1.081172	0.6897600	0.8820878
8.00	1.085143	0.6878698	0.8848315
16.00	1.085143	0.6878698	0.8848315

32.00	1.085143	0.6878698	0.8848315
64.00	1.085143	0.6878698	0.8848315
128.00	1.085143	0.6878698	0.8848315
256.00	1.085143	0.6878698	0.8848315
512.00	1.085143	0.6878698	0.8848315
1024.00	1.085143	0.6878698	0.8848315
2048.00	1.085143	0.6878698	0.8848315

Tuning parameter 'sigma' was held constant at a value of 0.01632049
 RMSE was used to select the optimal model using the smallest value.
 The final values used for the model were sigma = 0.01632049 and C = 2.

```
svm_Pred2 <- predict(svm_model2, newdata = df_test)

## The function 'postResample' can be used to get the test set performance values
svm_pv2 <- postResample(pred = svm_Pred2, obs = df_test$Yield)

# Model performance metrics
sum_t2 <- data.frame(
  knn_pv2,
  nnet_pv2,
  avnnet_pv2,
  mars_pv2,
  svm_pv2
)

print(sum_t2)
```

	knn_pv2	nnet_pv2	avnnet_pv2	mars_pv2	svm_pv2
RMSE	1.4986649	1.6189836	1.6190778	1.4131453	1.2756876
Rsquared	0.3302198	0.3159727	0.3159386	0.3911259	0.5011139
MAE	1.2035764	1.3311308	1.3312383	1.0797807	1.0200509


SVM regression gives the optimal performance based on RMSE, Rsquared and MAE on test set.

- (b) Which predictors are most important in the optimal nonlinear regression model?
 Do either the biological or process variables dominate the list? How do the top ten important predictors compare to the top ten predictors from the optimal linear model?

In linear model, ManufacturingProcess32 was the most important predictor but in non-linear model, it is 2nd most important predictor - the most important predictor is ManufacturingProcess13.

In linear model, only 2 of top 10 were Biological where as in non-linear, 4 of them were.

```
# code
varimp <- varImp(svm_model2, scale=F, useModel = T)
plot(varimp, top=15, scales = list(y = list(cex = 0.8)))
```




- (c) Explore the relationships between the top predictors and the response for the predictors that are unique to the optimal nonlinear regression model. Do these plots reveal intuition about the biological or process predictors and their relationship with yield?

From Bivariate plot and correlation matrix, we know that **ManufacturingProcess32** has fairly positive relationship with Yield where as other 2 variables have fairly negative relationship. Among biological predictors, we know **BiologicalMaterial06** is the most important with fairly strong positive relationship with Yield.

This information can help researchers to focus more on **ManufacturingProcess32** and **BiologicalMaterial06** if their goal is to increase Yield.

```
# code
viporder <- order(abs(varimp$importance),decreasing=TRUE)
topVIP <- rownames(varimp$importance)[viporder[c(1:5)]]

# bivariate relationship
featurePlot(df_train[, topVIP],
            df_train$Yield,
            plot = "scatter",
            between = list(x = 1, y = 1),
            type = c("g", "p", "smooth"),
            layout = c(5,1),
            labels = rep("", 2))
```



```
# corr_matrix
corr_top5 <- cor(df_train[, topVIP], df_train$Yield, method = 'pearson', use = 'pairwise.complete.obs')
corr_top5
```

[,1]

ManufacturingProcess13	-0.5645290
ManufacturingProcess32	0.6017957
ManufacturingProcess17	-0.4666781
BiologicalMaterial06	0.4673292
ManufacturingProcess09	0.5423679