

DATA 624: Project 1 - Part B

Sang Yoon (Andy) Hwang & Vinicio Haro

October 22, 2019

Contents

Part B: Forecasting Power	3
Exploration	3
Data Model	5
Forecast	7
Discussion	7
Appendix	9
Part B	9

Part B: Forecasting Power

Instructions: Part B consists of a simple dataset of residential power usage for January 1998 until December 2013. Your assignment is to model these data and a monthly forecast for 2014. The data is given in a single file. The variable 'KWH' is power consumption in Kilowatt hours, the rest is straight forward. Add these to your existing files above - clearly labeled.

Exploration

From our time series data (frequency = 12, monthly power_data) we observed there is a missing value in September 2008. We used imputation method called na.interpolation which performs a technique in numerical analysis which estimates a value from known data points. For our case, linear method using first order Taylor polynomial is used.

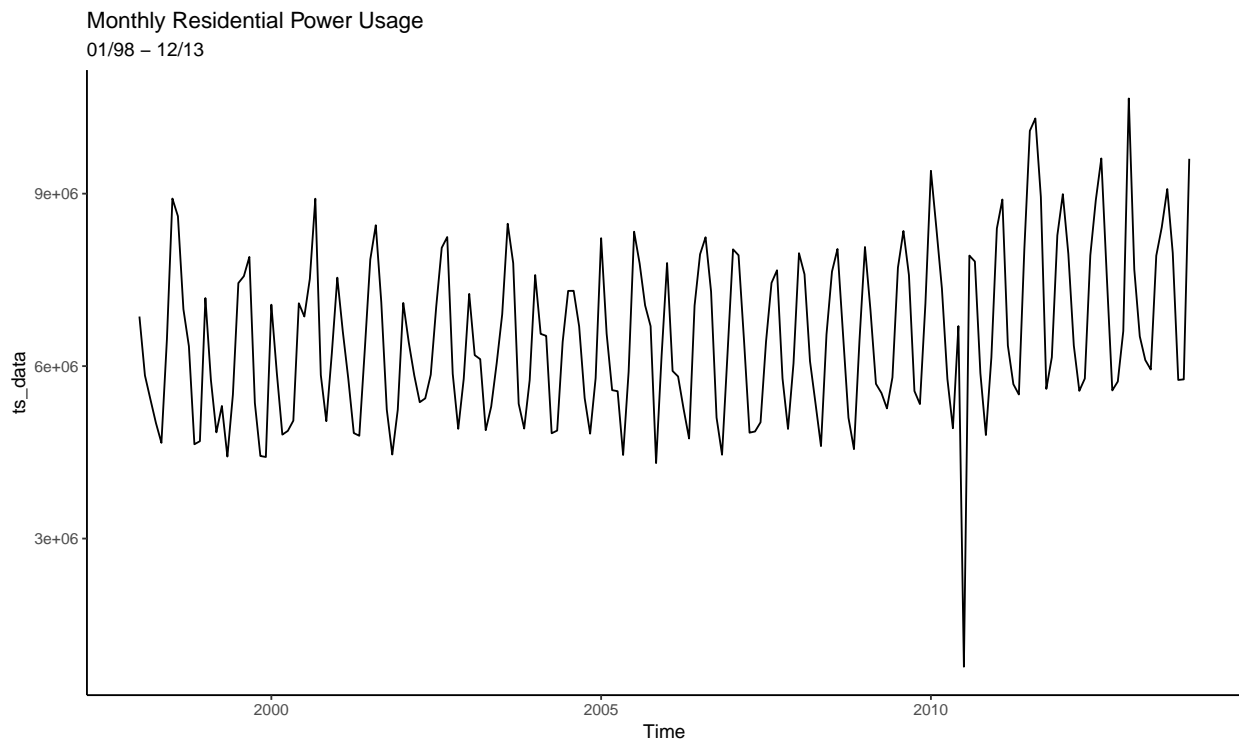
Our initial time series plot reveal annual seasonality within this time series. The box plot/seasonality plot actually reveals where power consumption fluctuations occur within each of the cycle positions. We can speculate that this could be due to there being no major Holidays that require power draining decor plus we assume minimal AC usage during the cold months.

We see power consumption increase between the months of June and August. This must be tied to AC usage during the warmer months of a year and finally power usage dips from September to November with a small spike in December. We speculate that this is due to transitioning out of summer. The spike in December could be connected to the usage of Holiday lights being kept on.

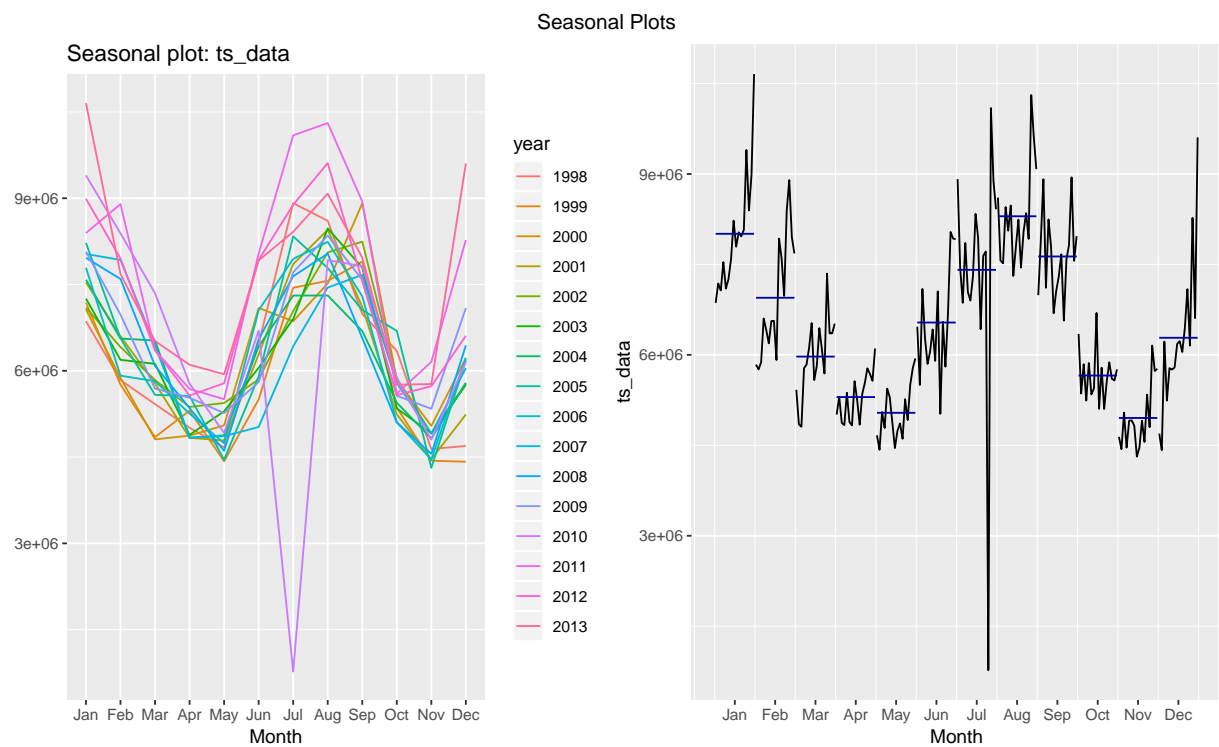
Within the overall TS plot, we see a dip in July 2010. This could be due to a power outage during a hot summer month. This can certainly be considered to be an outlier within this TS. Using TSO outliers, we can actually identify the index where our outliers may be. TSO outliers also replaces the outlier using Box-Cox. If set lambda=auto, then TSO outliers will automatically perform Box-Cox transformation.

The ACF plot shows that autocorrelations are well outside the significant space indicating the series is not white noise, non-stationary.

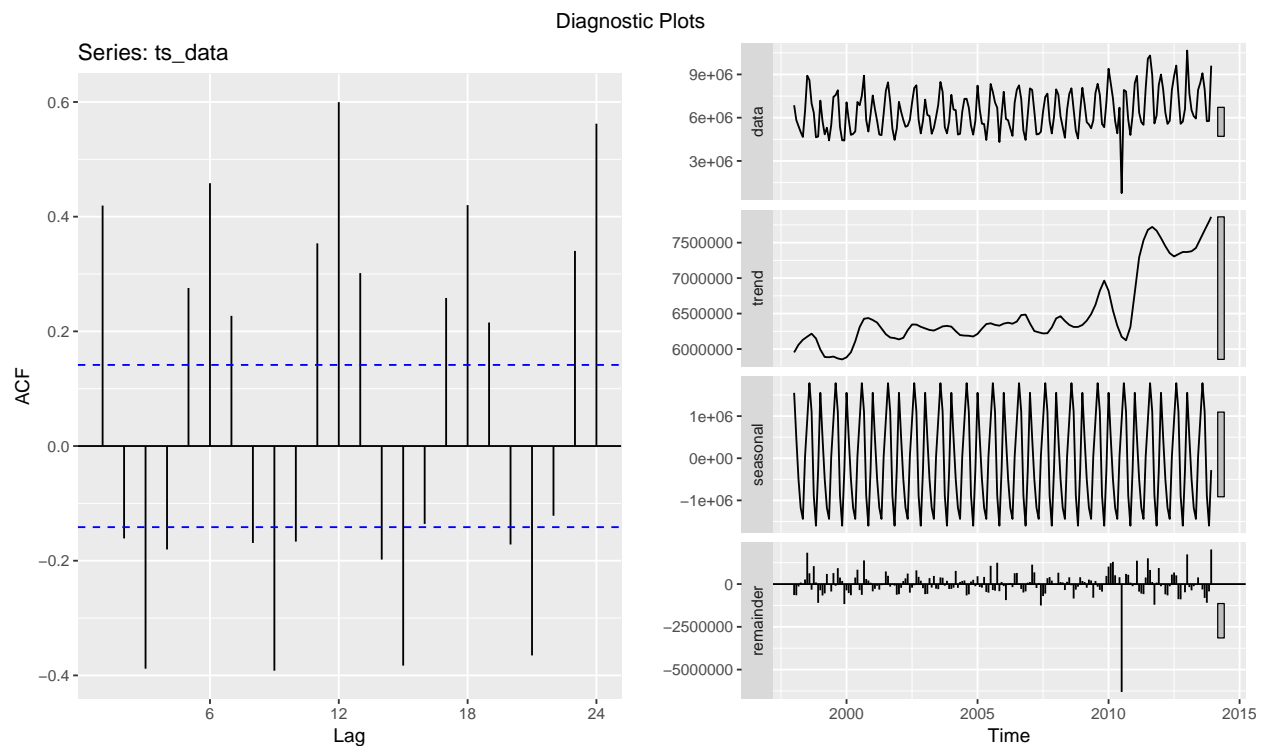
0.0.1 Series plot



0.0.2 Seasonal plots



0.0.3 Diagnostic plots



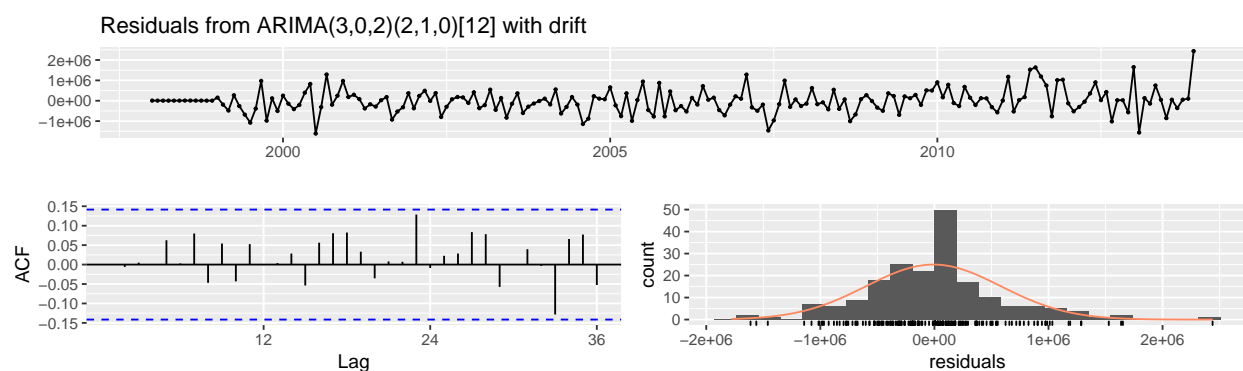
Data Model

Out of the models we built, we can make some preliminary observations. The residuals for each of our models does not have a major deviance from normality, however residuals of Model #1: ARIMA do not have an extended number of bins distorting the normality proximity but we can say it is still fairly normally distributed.

The residual ACF plots show residual autocorrelations for each of our models. Model #1: ARIMA has less autocorrelation than the other three models. Model 1 is well within the 95% limits indicated by the dotted blue lines.

If we examine the Ljung-Box test results for our models, the only model with a p-value > 0.05 is Model #1: ARIMA. This implies that the residuals from other models are not independent, hence not white noise. The full model summary can be viewed in the appendix.

0.0.4 Model #1: ARIMA

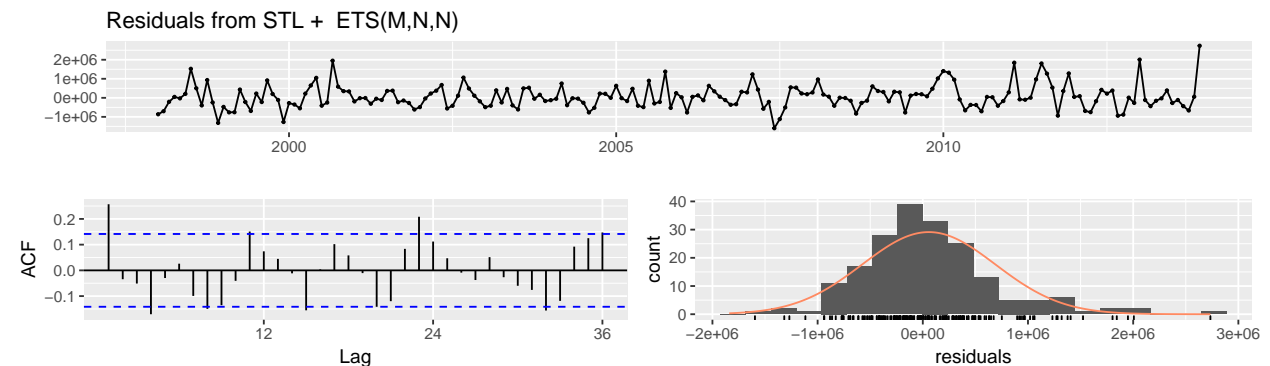


```

FALSE
FALSE  Ljung-Box test
FALSE
FALSE data:  Residuals from ARIMA(3,0,2)(2,1,0)[12] with drift
FALSE Q* = 12.555, df = 16, p-value = 0.705
FALSE
FALSE Model df: 8.    Total lags used: 24

```

0.0.5 Model #2: STL (no-demped) - MNN

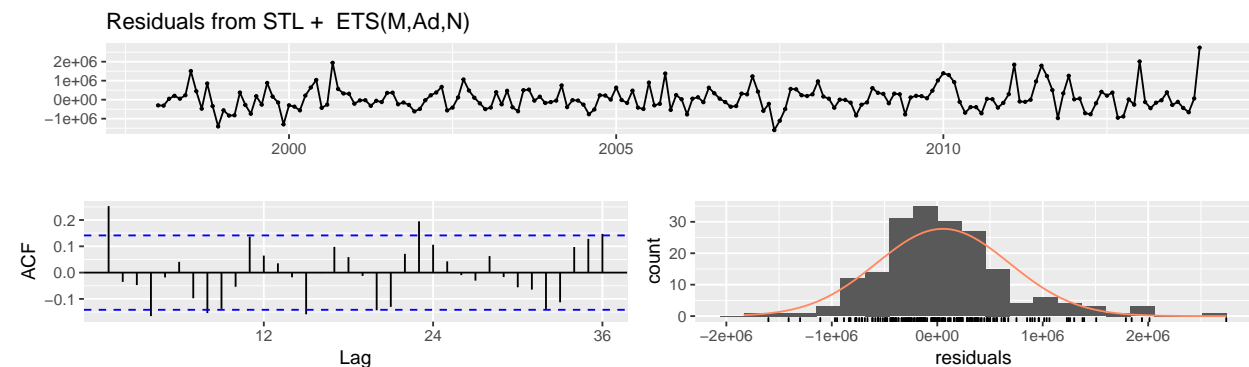


```

FALSE
FALSE  Ljung-Box test
FALSE
FALSE data:  Residuals from STL + ETS(M,N,N)
FALSE Q* = 65.934, df = 22, p-value = 2.84e-06
FALSE
FALSE Model df: 2.    Total lags used: 24

```

0.0.6 Model #2-2: STL (demped) - MAdN




```

FALSE
FALSE  Ljung-Box test
FALSE
FALSE data:  Residuals from STL + ETS(M,Ad,N)
FALSE Q* = 63.375, df = 19, p-value = 1.119e-06
FALSE
FALSE Model df: 5.    Total lags used: 24

```

0.0.7 Model #3: ets - MNM




```
FALSE
FALSE  Ljung-Box test
FALSE
FALSE data:  Residuals from ETS(M,N,M)
FALSE Q* = 32.042, df = 10, p-value = 0.000394
FALSE
FALSE Model df: 14.    Total lags used: 24
```

Forecast

`auto.arima()` performs cross validation on hyperparameter tuning to find the best model with parameters of order and seasonal that minimize AIC. This gave us **arima_model**: ARIMA(3, 0, 2)(2, 1, 0)12 with drift resulting AIC = 5332.24.

Since ARIMA is the only reliable model, as other models failed Ljung test, we will plot forecasts of ARIMA only. The forecasted values can be viewed in the appendix.

0.0.8 Model #1: ARIMA



Discussion

We implemented a cross validation method of testing for $h=12$. The process randomly chooses 12 points to measure and take the average of RMSEs. By definition, a lower RMSE on test set is attributed with a better forecast on unseen data.

Using Time series cross-validation, we compute RMSE on testset ($h=12$). We would have to pick the model with the lowest RMSE on test set as our final model if we had more than 1 model to compare. In our case, since we only have 1 model left after Ljung test, we have no choice but to pick seasonal ARIMA model as our final choice. Cross-validation test shows that RMSE on test is

around 720k when RMSE on training is around 589k. We can conclude the model is not necessarily overfitted. Given that MAPE on training is less than 7, it is not a surprising result.

```
FALSE [1] "RMSE - train: 589381.7"
```

```
FALSE [1] "RMSE - test: 725175"
```


Appendix

Part B

Model Summary

ARIMA:

FALSE

FALSE Forecast method: ARIMA(3,0,2)(2,1,0)[12] with drift

FALSE

FALSE Model Information:

FALSE Series: ts_data_o

FALSE ARIMA(3,0,2)(2,1,0)[12] with drift

FALSE

FALSE Coefficients:

	ar1	ar2	ar3	ma1	ma2	sar1	sar2	drift
--	-----	-----	-----	-----	-----	------	------	-------

	-0.5606	-0.2216	0.3284	0.8902	0.4827	-0.7249	-0.4152	9018.405
--	---------	---------	--------	--------	--------	---------	---------	----------

s.e.	0.3992	0.3382	0.0960	0.4120	0.4551	0.0797	0.0841	3027.685
------	--------	--------	--------	--------	--------	--------	--------	----------

FALSE

FALSE sigma^2 estimated as 3.878e+11: log likelihood=-2657.12

FALSE AIC=5332.24 AICc=5333.3 BIC=5360.97

FALSE

FALSE Error measures:

	ME	RMSE	MAE	MPE	MAPE	MASE
--	----	------	-----	-----	------	------

FALSE Training set	-8455.077	589381.7	427752.5	-0.7944782	6.475365	0.6904053
--------------------	-----------	----------	----------	------------	----------	-----------

FALSE ACF1

FALSE Training set 0.0006090194

FALSE

FALSE Forecasts:

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
--	----------------	-------	-------	-------	-------

FALSE Jan 2014	10210619	9412589	11008649	8990138	11431100
----------------	----------	---------	----------	---------	----------

FALSE Feb 2014	8722658	7882412	9562903	7437613	10007702
----------------	---------	---------	---------	---------	----------

FALSE Mar 2014	7137962	6295514	7980411	5849548	8426376
----------------	---------	---------	---------	---------	---------

FALSE Apr 2014	5919874	5060514	6779234	4605596	7234152
----------------	---------	---------	---------	---------	---------

FALSE May 2014	5946730	5087082	6806377	4632012	7261448
----------------	---------	---------	---------	---------	---------

FALSE Jun 2014	8383812	7524148	9243475	7069070	9698553
----------------	---------	---------	---------	---------	---------

FALSE Jul 2014	9362213	8500206	10224219	8043888	10680538
----------------	---------	---------	----------	---------	----------

FALSE Aug 2014	10018953	9155935	10881971	8699080	11338826
----------------	----------	---------	----------	---------	----------

FALSE Sep 2014	8547612	7684559	9410664	7227687	9867536
----------------	---------	---------	---------	---------	---------

FALSE Oct 2014	5781906	4918467	6645344	4461391	7102421
----------------	---------	---------	---------	---------	---------

FALSE Nov 2014	6193673	5329717	7057629	4872367	7514980
----------------	---------	---------	---------	---------	---------

FALSE Dec 2014	8373767	7509705	9237829	7052298	9695236
----------------	---------	---------	---------	---------	---------

STL - MNN:

FALSE

FALSE Forecast method: STL + ETS(M,N,N)

FALSE

FALSE Model Information:

FALSE ETS(M,N,N)

```

FALSE
FALSE Call:
FALSE ets(y = x, model = etsmodel, allow.multiplicative.trend = allow.multiplicative.trend)
FALSE
FALSE Smoothing parameters:
FALSE alpha = 0.1159
FALSE
FALSE Initial states:
FALSE l = 6317745.8917
FALSE
FALSE sigma: 0.097
FALSE
FALSE AIC AICc BIC
FALSE 6139.631 6139.758 6149.403
FALSE
FALSE Error measures:
FALSE ME RMSE MAE MPE MAPE MASE
FALSE Training set 56926.03 633571.7 460713.4 -0.03288687 6.945185 0.7436052
FALSE ACF1
FALSE Training set 0.2570241
FALSE
FALSE Forecasts:
FALSE Point Forecast Lo 80 Hi 80 Lo 95 Hi 95
FALSE Jan 2014 8992609 8049591 9935628 7550387 10434831
FALSE Feb 2014 7908116 6958724 8857508 6456146 9360086
FALSE Mar 2014 7079434 6123709 8035158 5617779 8541088
FALSE Apr 2014 6435209 5473193 7397225 4963933 7906486
FALSE May 2014 6161593 5193326 7129860 4680756 7642430
FALSE Jun 2014 7728705 6754226 8703185 6238368 9219043
FALSE Jul 2014 8837980 7857327 9818633 7338201 10337759
FALSE Aug 2014 9376841 8390053 10363630 7867678 10886004
FALSE Sep 2014 8601001 7608114 9593888 7082511 10119490
FALSE Oct 2014 6670419 5671470 7669368 5142658 8198180
FALSE Nov 2014 6035845 5030870 7040821 4498868 7572822
FALSE Dec 2014 7189087 6178120 8200053 5642947 8735226

STL - MadN:

FALSE
FALSE Forecast method: STL + ETS(M,Ad,N)
FALSE
FALSE Model Information:
FALSE ETS(M,Ad,N)
FALSE
FALSE Call:
FALSE ets(y = x, model = etsmodel, damped = TRUE, allow.multiplicative.trend = allow.multiplicative.tr
FALSE
FALSE Smoothing parameters:
FALSE alpha = 0.1233
FALSE beta = 1e-04
FALSE phi = 0.8
FALSE
FALSE Initial states:
FALSE l = 5615471.7851
FALSE b = 173606.4508

```

```

FALSE
FALSE    sigma:    0.0972
FALSE
FALSE      AIC      AICc      BIC
FALSE 6143.452 6143.906 6162.997
FALSE
FALSE Error measures:
FALSE              ME      RMSE      MAE      MPE      MAPE      MASE
FALSE Training set 54337.68 631081.9 458777.5 -0.07364717 6.937249 0.7404807
FALSE              ACF1
FALSE Training set 0.2528558
FALSE
FALSE Forecasts:
FALSE      Point Forecast    Lo 80      Hi 80      Lo 95      Hi 95
FALSE Jan 2014      9007707 8060947 9954467 7559763 10455651
FALSE Feb 2014      7923348 6969325 8877372 6464295 9382401
FALSE Mar 2014      7094774 6133536 8056011 5624687 8564860
FALSE Apr 2014      6450635 5482232 7419038 4969591 7931680
FALSE May 2014      6177088 5201569 7152607 4685160 7669016
FALSE Jun 2014      7744256 6761668 8726843 6241518 9246993
FALSE Jul 2014      8853574 7863967 9843182 7340100 10367048
FALSE Aug 2014      9392471 8395890 10389052 7868332 10916609
FALSE Sep 2014      8616658 7613151 9620166 7081926 10151391
FALSE Oct 2014      6686100 5675711 7696488 5140843 8231356
FALSE Nov 2014      6051544 5034319 7068769 4495832 7607255
FALSE Dec 2014      7204799 6180782 8228817 5638700 8770899

ets - MNM:

FALSE
FALSE Forecast method: ETS(M,N,M)
FALSE
FALSE Model Information:
FALSE ETS(M,N,M)
FALSE
FALSE Call:
FALSE ets(y = ts_data_o)
FALSE
FALSE Smoothing parameters:
FALSE     alpha = 0.1428
FALSE     gamma = 0.2119
FALSE
FALSE Initial states:
FALSE     l = 6189149.8743
FALSE     s = 0.8984 0.7596 0.938 1.2229 1.2597 1.2396
FALSE           1.0059 0.7638 0.8078 0.8864 1.0269 1.191
FALSE
FALSE    sigma:    0.0967
FALSE
FALSE      AIC      AICc      BIC
FALSE 6144.033 6146.760 6192.895
FALSE
FALSE Error measures:
FALSE              ME      RMSE      MAE      MPE      MAPE      MASE
FALSE Training set 45241.77 628252.5 481520.9 -0.04000239 7.277118 0.7771892

```

```

FALSE                               ACF1
FALSE Training set 0.1927438
FALSE
FALSE Forecasts:
FALSE      Point Forecast   Lo 80    Hi 80    Lo 95    Hi 95
FALSE Jan 2014      9917654 8689211 11146096 8038913 11796394
FALSE Feb 2014      8522973 7456477 9589469 6891908 10154038
FALSE Mar 2014      7012478 6126191 7898765 5657019 8367937
FALSE Apr 2014      6208601 5416196 7001006 4996722 7420480
FALSE May 2014      5928833 5164834 6692832 4760398 7097269
FALSE Jun 2014      7840532 6820624 8860440 6280717 9400347
FALSE Jul 2014      9115823 7919004 10312642 7285446 10946200
FALSE Aug 2014      9648549 8370229 10926869 7693527 11603571
FALSE Sep 2014      8553364 7409986 9696742 6804718 10302010
FALSE Oct 2014      6266745 5421655 7111835 4974291 7559199
FALSE Nov 2014      5938289 5130560 6746017 4702975 7173603
FALSE Dec 2014      8020901 6920610 9121192 6338151 9703651

```

R Script

```
# Dependencies
## processing
library(readxl)
library(tinytex)
library(readr)

## graphs
library(ggplot2)
library(janitor)
library(gridExtra)
library(grid)

## formatting
library(default)
library(knitr)
library(kableExtra)
library(tidyverse)
library(scales)
library(readxl)
library(lubridate)

## forecasting packages
library(fpp2)
library(forecast)

## outlier & imputation
library(imputeTS)
library(tsoutliers)

# load data
power_data <- read_csv("https://raw.githubusercontent.com/vindication09/DATA-624/master/ResidentialCust")

# Time Series
ts_data <- ts(power_data$KWH, frequency = 12, start = c(1998,1))

# Missing value imputation
ts_data <- na_interpolation(ts_data)

# STL decomposition
stl1 <- stl(ts_data, s.window = 'periodic')

# Handling outlier
outlier_func <- tsoutliers(ts_data, iterate = 2, lambda = "auto")

# Time Series - After outlier and imputation handled
ts_data_o <- ts_data # Let's treat outlier handled data separately for Modelling part.
ts_data_o[outlier_func$index] <- outlier_func$replacements

# Model#1: ARIMA
arma_auto <- auto.arima(ts_data_o)
arma_model <- forecast(arma_auto, h=12)
```

```

# Model #2: STL (no-demped) - MNN
stl_ndemp <- stlf(ts_data_o, s.window = "periodic", robust=TRUE, h = 12)

# Model #2-2: STL (demped) - MAdN
stl_demp <- stlf(ts_data_o, damped=TRUE, s.window = "periodic", robust=TRUE, h = 12)

# Model #3: ets - MNM
ets_auto <- ets(ts_data_o)
ets_model <- forecast(ets_auto, h=12)

# tsCv - ARIMA -> it takes so much time. I got the results and saved them
##arima_cv <- function(x, h){forecast(Arima(x, order = c(3, 0, 2), seasonal = c(2, 1, 0), include.drift
##e <- tsCV(ts_data_o, arima_cv, h=12)

# RMSEs -> tsCV takes lot of time to process so just saved the output
#rmse_train_arima <- arima_auto[2]
#rmse_test_arima <- sqrt(mean(e^2, na.rm=TRUE))

rmse_train_arima <- 589381.7
rmse_test_arima <- 725175

# Save output
write.csv(arima_model, file="forecasts/POWER_ARIMA_FC.csv")

```